


Neural Network Example

Build a 2-hidden layers fully connected neural network (a.k.a multilayer perceptron) with TensorFlow v2.

This example is using a low-level approach to better understand all mechanics behind building neural networks and the training process.

- Author: Aymeric Damien
- Project: <https://github.com/aymericdamien/TensorFlow-Examples/>

Neural Network Overview



MNIST Dataset Overview

This example is using MNIST handwritten digits. The dataset contains 60,000 examples for training and 10,000 examples for testing. The digits have been size-normalized and centered in a fixed-size image (28x28 pixels) with values from 0 to 255.

In this example, each image will be converted to float32, normalized to $[0, 1]$ and flattened to a 1-D array of 784 features (28×28).

 MNIST Dataset

More info: <http://yann.lecun.com/exdb/mnist/>

```
1 from __future__ import absolute_import, division, print_function
2 from sklearn import preprocessing
3 import tensorflow as tf
4 from tensorflow.keras import Model, layers
5 import numpy as np
6 import pandas as pd          # For loading and processing the dataset.
7 from sklearn.model_selection import train_test_split

1 # MNIST dataset parameters.
2 num_classes = 2 # total classes (0-9 digits).
3 num_features = 14 # data features (img shape: 28*28).
4
5 # Training parameters.
6 learning_rate = 0.1
7 training_steps = 2000
8 batch_size = 256
9 display_step = 100
10
11 # Network parameters.
12 n_hidden_1 = 128 # 1st layer number of neurons.
13 n_hidden_2 = 256 # 2nd layer number of neurons.

1 # Finally, we convert the Pandas dataframe to a NumPy array, and split it into a t
2 # Read the CSV input file and show first 5 rows
3 df_train = pd.read_csv('/content/task1_dataset_train.csv')
4 df_train.head(5)
5 X = df_train.drop('y', axis=1).values
6 y = df_train['y'].values
7 min_max_scaler = preprocessing.MinMaxScaler()
8 X_scale = min_max_scaler.fit_transform(X)
9 X_train, X_test, y_train, y_test = train_test_split(X_scale, y, test_size=0.2)

1

1 # Use tf.data API to shuffle and batch data.
2 train_data = tf.data.Dataset.from_tensor_slices((X_train, y_train))
3 train_data = train_data.repeat().shuffle(5000).batch(batch_size).prefetch(1)

1

1 # Create TF Model.
2 class NeuralNet(Model):
3     # Set layers.
4     def __init__(self):
5         super(NeuralNet, self).__init__()
6         # First fully-connected hidden layer.
7         self.fc1 = layers.Dense(n_hidden_1, activation=tf.nn.relu)
```

```

8         # First fully-connected hidden layer.
9         self.fc2 = layers.Dense(n_hidden_2, activation=tf.nn.relu)
10        # Second fully-connected hidden layer.
11        self.out = layers.Dense(num_classes)
12
13        # Set forward pass.
14        def call(self, x, is_training=False):
15            x = self.fc1(x)
16            x = self.fc2(x)
17            x = self.out(x)
18            if not is_training:
19                # tf cross entropy expect logits without softmax, so only
20                # apply softmax when not training.
21                x = tf.nn.softmax(x)
22            return x
23
24        # Build neural network model.
25        neural_net = NeuralNet()

1 # Cross-Entropy Loss.
2 # Note that this will apply 'softmax' to the logits.
3 def cross_entropy_loss(x, y):
4     # Convert labels to int 64 for tf cross-entropy function.
5     y = tf.cast(y, tf.int64)
6     # Apply softmax to logits and compute cross-entropy.
7     loss = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y, logits=x)
8     # Average loss across the batch.
9     return tf.reduce_mean(loss)
10
11 # Accuracy metric.
12 def accuracy(y_pred, y_true):
13     # Predicted class is the index of highest score in prediction vector (i.e. arg
14     correct_prediction = tf.equal(tf.argmax(y_pred, 1), tf.cast(y_true, tf.int64))
15     return tf.reduce_mean(tf.cast(correct_prediction, tf.float32), axis=-1)
16
17 # Stochastic gradient descent optimizer.
18 optimizer = tf.optimizers.SGD(learning_rate)

1 # Optimization process.
2 def run_optimization(x, y):
3     # Wrap computation inside a GradientTape for automatic differentiation.
4     with tf.GradientTape() as g:
5         # Forward pass.
6         pred = neural_net(x, is_training=True)
7         # Compute loss.
8         loss = cross_entropy_loss(pred, y)
9
10    # Variables to update, i.e. trainable variables.
11    trainable_variables = neural_net.trainable_variables
12

```

```
13 # Compute gradients.
14 gradients = g.gradient(loss, trainable_variables)
15
16 # Update W and b following gradients.
17 optimizer.apply_gradients(zip(gradients, trainable_variables))

1 # Run training for the given number of steps.
2 for step, (batch_x, batch_y) in enumerate(train_data.take(training_steps), 1):
3     # Run the optimization to update W and b values.
4     run_optimization(batch_x, batch_y)
5
6     if step % display_step == 0:
7         pred = neural_net(batch_x, is_training=True)
8         loss = cross_entropy_loss(pred, batch_y)
9         acc = accuracy(pred, batch_y)
10        print("step: %i, loss: %f, accuracy: %f" % (step, loss, acc))
```

```
step: 100, loss: 0.486165, accuracy: 0.769531
step: 200, loss: 0.520279, accuracy: 0.750000
step: 300, loss: 0.530207, accuracy: 0.714844
step: 400, loss: 0.465484, accuracy: 0.789062
step: 500, loss: 0.461355, accuracy: 0.792969
step: 600, loss: 0.475713, accuracy: 0.789062
step: 700, loss: 0.436438, accuracy: 0.804688
step: 800, loss: 0.450595, accuracy: 0.757812
step: 900, loss: 0.489487, accuracy: 0.765625
step: 1000, loss: 0.461266, accuracy: 0.750000
step: 1100, loss: 0.420362, accuracy: 0.804688
step: 1200, loss: 0.434839, accuracy: 0.812500
step: 1300, loss: 0.410994, accuracy: 0.781250
step: 1400, loss: 0.398383, accuracy: 0.800781
step: 1500, loss: 0.347369, accuracy: 0.851562
step: 1600, loss: 0.385559, accuracy: 0.796875
step: 1700, loss: 0.425821, accuracy: 0.789062
step: 1800, loss: 0.377654, accuracy: 0.824219
step: 1900, loss: 0.424663, accuracy: 0.769531
step: 2000, loss: 0.415788, accuracy: 0.796875
```

```
1 # Test model on validation set.
2 pred = neural_net(X_test, is_training=False)
3 print("Test Accuracy: %f" % accuracy(pred, y_test))
```

```
Test Accuracy: 0.807301
```

```
1 df_test = pd.read_csv('/content/task1_dataset_test.csv')
2 df_test.head()
```

	x_1	x_2	x_3	x_4	x_5	x_6	x_7	x_8	x_9	x_10	x_11	x_12	x_13	x_14
0	57	2	21598	1	1	6	0	0	0	1	1	0	19	0
1	27	2	969	0	0	0	9	0	0	1	1	0	19	0
2	25	6	1126	1	1	1	6	1	0	0	12	0	4	0

```

1 df_test = df_test.values
2 min_max_scaler = preprocessing.MinMaxScaler()
3 df_test_scale = min_max_scaler.fit_transform(df_test)

```

```

1 X_test

```

```

array([[0.27777778, 0.25          , 0.15548456, ..., 0.          , 0.          ,
        0.          ],
       [0.125        , 0.25          , 0.12631384, ..., 0.03296703, 0.04301075,
        0.          ],
       [0.79166667, 0.625          , 0.11913692, ..., 0.          , 0.29032258,
        0.          ],
       ...,
       [0.55555556, 0.625          , 0.61994721, ..., 0.          , 0.          ,
        0.          ],
       [0.06944444, 0.125          , 0.20924202, ..., 0.          , 0.04301075,
        0.          ],
       [0.40277778, 0.25          , 0.69912488, ..., 0.          , 0.          ,
        0.          ]])

```

```

1 # Predict 5 images from validation set.
2 n_t = 100
3 test_t = df_test_scale[:n_t]
4 predictions = neural_net(test_t)
5
6 # Display image and model prediction.
7 for i in range(n_t):
8     print("Model prediction: %i" % np.argmax(predictions.numpy()[i]))

```

```

☐ Model prediction: 0
Model prediction: 0
Model prediction: 1
Model prediction: 0
Model prediction: 1
Model prediction: 0
Model prediction: 1
Model prediction: 0
Model prediction: 0
Model prediction: 0
Model prediction: 0
Model prediction: 0
Model prediction: 0
Model prediction: 0
Model prediction: 0
Model prediction: 0
Model prediction: 0

```

License: MIT

✓ 0s completed at 12:09 PM

● ✕